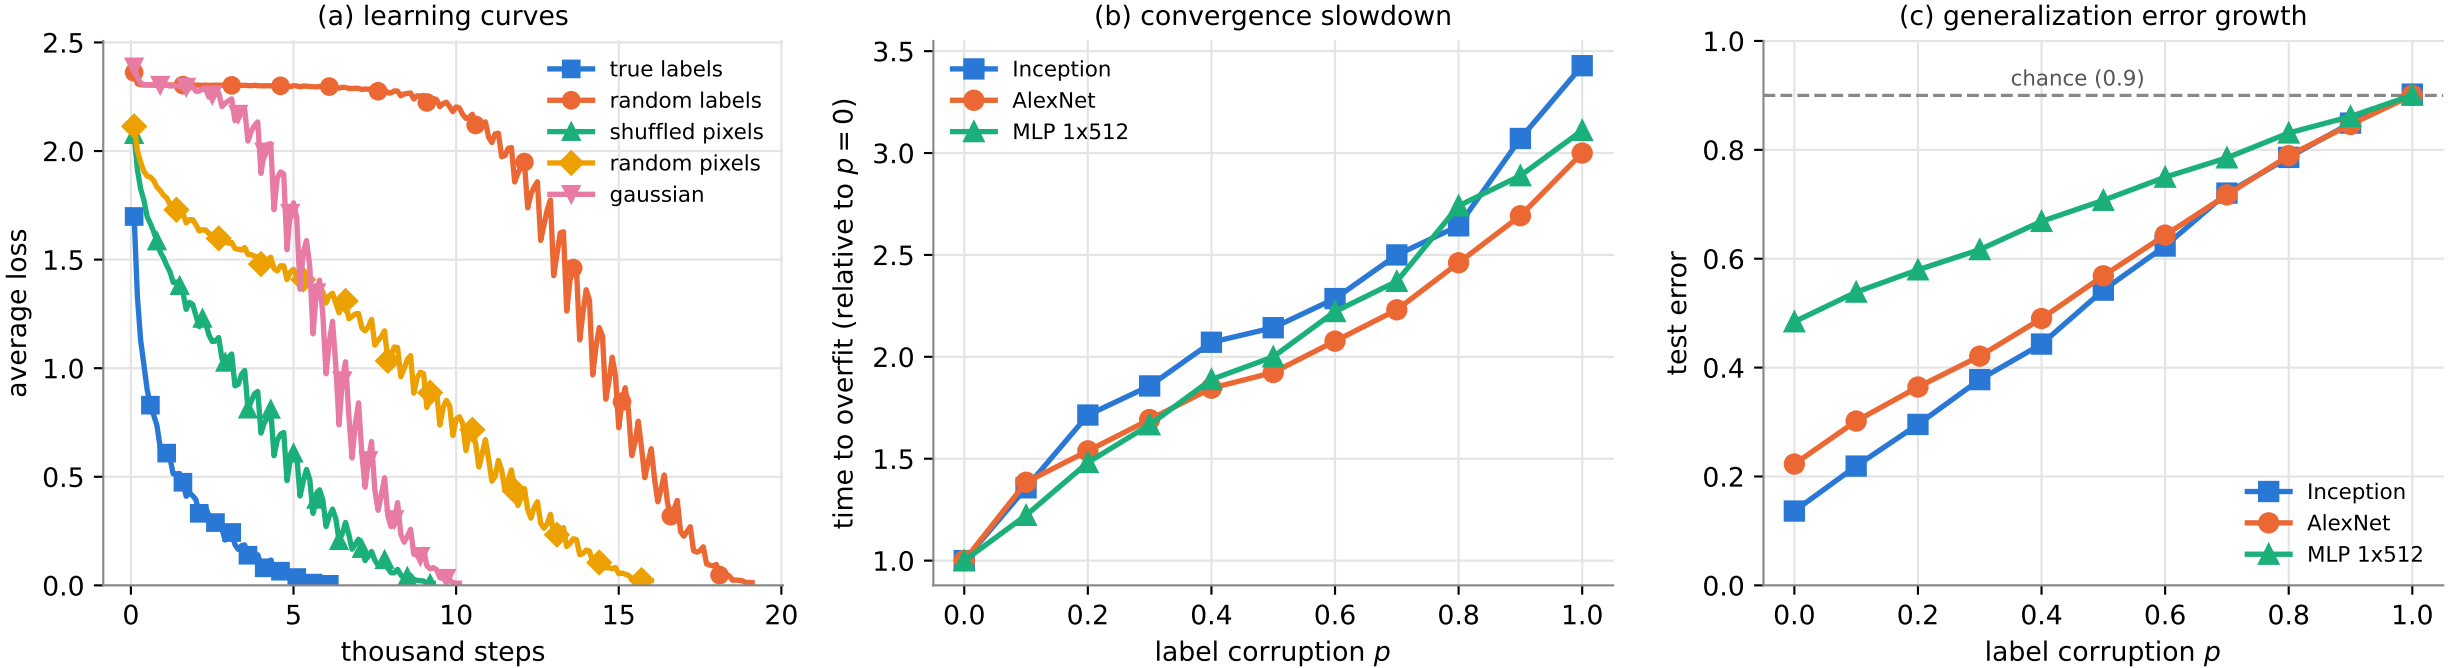
(a) learning curves
average loss
thousand steps
true labels
random labels
shuffled pixels
random pixels
gaussian
(b) convergence slowdown
time to overfit (relative to $p = 0$)
label corruption p
Inception
AlexNet
MLP 1x512
(c) generalization error growth
test error
label corruption p
chance (0.9)
Inception
AlexNet
MLP 1x512